

MEMORY-EFFICIENT INFERENCE FOR LARGE LANGUAGE MODELS

CROSS-REFERENCE TO RELATED APPLICATIONS

This application claims the benefit of U.S. Provisional Application No. 63/000,000, filed January 5, 2025, which is incorporated herein by reference.

BACKGROUND

Modern systems require efficient inference.

[INSERT DESCRIPTION OF PRIOR ART]

BRIEF DESCRIPTION OF THE DRAWINGS

FIG. 1 illustrates a system overview.

FIG. 2 shows a method flow.

DETAILED DESCRIPTION

Referring to FIG. 1, the system 100 includes a processor 110 configured to perform memory-efficient inference for large language models.

What is claimed is:

1. A method comprising: receiving input data; and processing the input data with a processor to produce an inference result.
2. The method of claim 1, further comprising storing the inference result.
3. The method of claim 1, wherein processing comprises quantized inference.

ABSTRACT

A method and system for memory-efficient inference for large language models, including a processor configured to process input data and produce an inference result.