

MEMORY-EFFICIENT INFERENCE FOR LARGE LANGUAGE MODELS

BACKGROUND

Modern systems require efficient inference.