

MEMORY-EFFICIENT INFERENCE FOR LARGE LANGUAGE MODELS

BACKGROUND

Modern systems require efficient inference.

[INSERT DESCRIPTION OF PRIOR ART]